

✓ Credit Risk Classification Using Machine Learning

CIND 820 – Capstone Project

Strahinja Nakic | Toronto Metropolitan University | 500809487

✓ Project Overview

This project uses the German Credit dataset to classify loan applicants as good or bad credit risks. The analysis applies three supervised machine learning models: Logistic Regression, Decision Tree, and Naive Bayes. These models are trained and evaluated using performance metrics such as accuracy, recall, confusion matrix, and cross-validation

```
## Importing the Libraries
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, recall_score, confusion_matrix, classification_report
import pprint
from google.colab import drive
```

✓ Mounting Google Drive and Loading Dataset

The dataset is stored in Google Drive and loaded using pandas. Column names are manually defined since the dataset has no header row

```
## Mounting my Google Drive
drive.mount('/content/drive')

## Loading the dataset from Google Drive
file_path = '/content/drive/My Drive/CIND 820/german.data' ## Replace with your d

column_names = [
    "checking_account", "duration", "credit_history", "purpose", "credit_amount",
    "savings_account", "employment", "installment_rate", "personal_status",
    "other_debtors", "residence_since", "property", "age", "other_installment_plan",
    "housing", "existing_credits", "job", "people_liable", "telephone",
    "foreign_worker", "credit_risk"
]
```

```
1
```

```
df = pd.read_csv(file_path, sep=' ', header=None, names=column_names)
```

✓ Preprocessing

I am converting the target variable into binary (1 = good credit, 0 = bad credit). Then I will label encode all categorical columns and scale the features using StandardScaler

```
## Converting the target to binary
df['credit_risk'] = df['credit_risk'].map({1: 1, 2: 0})

## Dropping rows with NaN values in the credit_risk column
df.dropna(subset=['credit_risk'], inplace=True)

## Label encoding all text columns
label_encoder = LabelEncoder()
for column in df.columns:
    if df[column].dtype == 'object':
        df[column] = label_encoder.fit_transform(df[column])

## Separating features and target
X = df.drop("credit_risk", axis=1)
y = df["credit_risk"]

## Scaling the features for better model performance and efficiency
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

✓ Splitting my Data

I am using 80 percent for training and 20 percent for testing. Stratify is used to maintain class balance.

```
## Splitting into train and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=1, stratify=y
)
```

✓ Defining and Training the Models Explanation

The three classification models I will train are Logistic Regression, Decision Tree, Naive Bayes

```
## Defining and training my models
models = {
    "Logistic Regression": LogisticRegression(max_iter=3000),
    "Decision Tree": DecisionTreeClassifier(random_state=1),
    "Naive Bayes": GaussianNB()
}
```

✓ 6. Evaluating the Models

Now I will evaluate the models by using Accuracy, Recall, Confusion Matrix, 5-fold Cross-Validation Accuracy

```
## Evaluating my models
results = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)

    acc = accuracy_score(y_test, predictions)
    rec = recall_score(y_test, predictions)
    cm = confusion_matrix(y_test, predictions)
    report = classification_report(y_test, predictions, output_dict=True)
    cv_score = cross_val_score(model, X_train, y_train, cv=5).mean()

    results[name] = {
        "Accuracy": round(acc, 3),
        "Recall": round(rec, 3),
        "Cross-Validation Accuracy": round(cv_score, 3),
        "Confusion Matrix": cm
    }
```

✓ The Final Results

Below is a summary of the performance of each model. These results will be used in the final report to select the most appropriate model

```
## Showing the results
pprint.pprint(results)

{'Decision Tree': {'Accuracy': 0.72,
                   'Confusion Matrix': array([[ 35,  25],
                   [ 31, 109]]),
                   'Cross-Validation Accuracy': np.float64(0.684),
```

```
        'Recall': 0.779},  
'Logistic Regression': {'Accuracy': 0.76,  
                        'Confusion Matrix': array([[ 31,  29],  
            [ 19, 121]]),  
                        'Cross-Validation Accuracy': np.float64(0.76),  
                        'Recall': 0.864},  
'Naive Bayes': {'Accuracy': 0.76,  
                'Confusion Matrix': array([[ 39,  21],  
            [ 27, 113]]),  
                'Cross-Validation Accuracy': np.float64(0.721),  
                'Recall': 0.807}}
```